

Introduction to R

Module 2: Data Preparation using the Tidyverse

Andrew Proctor

andrew.proctor@phdstudent.hhs.se

February 5, 2018



- ① Intro
- ② Packages in R
- ③ Data Prep Preliminaries
- ④ Data Preparation
- ⑤ Cleaning data
- ⑥ Tidy data



Intro



Goals for Today

- ① Learn how to use packages in R
- ② Learn how to import and export data.
- ③ Learn how to perform common data prep functions from the tidyverse collection of packages.
- ④ Learn how to clean and “tidy” data.



Packages in R



Role of Packages in R

- Packages in R are similar to user-written commands (think *ssc install*) in Stata.
- But *most* things you do in Stata probably use core Stata commands.
- In R, most of your analysis will probably be done using packages.



CRAN
[Mirrors](#)
[What's new?](#)
[Task Views](#)
[Search](#)

About R
[R Homepage](#)

Available CRAN Packages By Date

Date	Package	
2018-02-03	fuzzyforest	Fuzzy Forests
2018-02-03	tuber	Client for the YouTube API
2018-02-02	adeqenet	Exploratory Analysis of Genetic and Gen
2018-02-02	antitrust	Tools for Antitrust Practitioners
2018-02-02	arsenal	An Arsenal of 'R' Functions for Large-S

Installing and using a package

- To install a package, use the function (preferably in the console) `install.packages()`
- To begin with, let's install 2 packages:
 - `tidyverse`: the umbrella package for common data preparation and visualization in R.
 - `rio`: a package for easy data import, export (saving), and conversion.

```
install.packages("tidyverse")    # Install tidyverse  
install.packages("rio")         # Install rio
```



Loading a package during analysis

Unlike Stata, in R you need to declare what packages you will be using at the beginning of each R document.

To do this, use the `library()` function. `__require()` also works, but it's use is discouraged for this purpose.

```
library("tidyverse")    # Install tidyverse  
library("rio")          # Install rio
```



Data Prep Preliminaries



Import and export using rio

Previously, importing and exporting data was a mess, with a lot of different functions for different file formats:

- Stata DTA files alone required two functions: `read.dta` (for Stata 6-12 DTA files), `read.dta13` (for Stata 13 and later files), etc.

The `rio` package simplifies this by reducing all of this to just one function, `import()`

- Automatically determines the file format of the file and uses the appropriate function from other packages to load in a file



Import and export using rio II

```
PISA_2015 <- import("PISA2015.sas7bdat")  
PISA_2015[1:5,1:6]
```

##	CNTRYID	CNT	CNTSCHID	CYC	NatCen	Region
## 1	8	ALB	800001	06MS	000800	800
## 2	8	ALB	800002	06MS	000800	800
## 3	8	ALB	800003	06MS	000800	800
## 4	8	ALB	800004	06MS	000800	800
## 5	8	ALB	800005	06MS	000800	800

```
export(PISA_2015, "PISA_2015.rds")
```



Tibbles: an update to the data frame

Last class, we covered data frames—the most basic data object class for data sets with a mix of data class.

Today, we introduce one final data object: the **tibble**!

The tibble can be thought of as an update to the data frame—and it's the first part of the *tidyverse* package that we'll look at.



Tibble vs data frames

There are three main benefits to the tibble:

① Displaying data frames:

- If you display a data frame, it will print as much as much output as allowed by the “max.print” option in the R environment. With large data sets, that’s far too much. Tibbles by default print the first 10 rows and as many columns as will fit in the window.

② Partial matching in data frames:

- When using the **\$** method to reference columns of a data frame, partial names will be matched if the reference isn’t exact. This might sound good, but the only real reason for there to be a partial match is a typo, in which case the match might be wrong.

③ Tibbles are required for some functions.



Creating or converting to tibbles

The syntax for creating tibbles exactly parallels the syntax for data frames:

- `tibble()` creates a tibble from underlying data or vectors.
- `as.tibble()` coerces an existing data object into a tibble.

```
PISA_2015 <- as.tibble(PISA_2015); PISA_2015[1:5,1:5]
```

```
## # A tibble: 5 x 5
##   CNTRYID CNT   CNTSCHID CYC   NatCen
##   <dbl> <chr>   <dbl> <chr> <chr>
## 1    8.00 ALB     800001 06MS  000800
## 2    8.00 ALB     800002 06MS  000800
## 3    8.00 ALB     800003 06MS  000800
## 4    8.00 ALB     800004 06MS  000800
## 5    8.00 ALB     800005 06MS  000800
```



Glimpse

Another tidyverse function that's very useful is `glimpse()`, a function very similar to `str()`.

- Both functions display information about the structure of a data object.
- `str()` provides more information, such as column (variable) attributes embedded from external data formats, but consequently is much less readable for complex data objects.
- `glimpse()` provides only column names, classes, and some data values (much more readable)
- I will often use `str()` when I want more detailed information about data structure, but use `glimpse()` for quicker glances at the data.



Pipes

Another major convenience enhancement from the tidyverse is ***pipes***, denoted **%>%**,

- Pipes allow you to combine multiple steps into a single piece of code.
- Specifically, after performing a function in one step, a pipe takes the data generated from the first step and uses it as the data input to a second step.



Pipes Example

```
barro.lee.data <- import("BL2013_MF1599_v2.1.dta") %>%  
  as.tibble() %>% glimpse(width = 50)
```

```
## Observations: 1,898  
## Variables: 20  
## $ BLcode      <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1...  
## $ country     <chr> "Algeria", "Algeria", "Al...  
## $ year        <dbl> 1950, 1955, 1960, 1965, 1...  
## $ sex         <chr> "MF", "MF", "MF", "MF", "...  
## $ agefrom     <dbl> 15, 15, 15, 15, 15, 15, 1...  
## $ age to      <dbl> 999, 999, 999, 999, 999, ...  
## $ lu          <dbl> 80.68459, 81.05096, 82.61...  
## $ lp          <dbl> 17.563400, 17.018442, 14....  
## $ lpc         <dbl> 3.745905, 3.464397, 3.069...  
## $ ls          <dbl> 1.454129, 1.639253, 2.752...
```



Data Preparation



Tidyverse and the verbs of data manipulation

A motivating principle behind the creation of the tidyverse was the language of programming should really behave like a language.

Data manipulation in the tidyverse is oriented around a few key “verbs” that perform common types of data manipulation.

- 1 `filter()` subsets the rows of a data frame based on their values.
- 2 `select()` selects variables (columns) based on their names.
- 3 `mutate()` adds new variables that are functions of existing variables.
- 4 `summarize()` creates a number of summary statistics out of many values.
- 5 `arrange()` changes the ordering of the rows.

Note: the first argument for each these functions is the data object (so pipe!).



Filtering data

Filtering keeps observations (rows) based on conditions.

- Just like using subset conditions in the row arguments of a bracketed subset

```
# Using brackets
```

```
wages[(wages$schooling > 10) & (wages$exper > 10),]
```

```
##           wage schooling      sex exper
## 2 249.6774          13 female      11
```

```
# Using filter
```

```
wages %>% filter(schooling > 10, exper > 10)
```

```
##           wage schooling      sex exper
## 1 249.6774          13 female      11
```



Filtering data ctd

Notice a couple of things about the output:

- 1 It doesn't look like we told `filter()` what data set we would be filtering.
 - That's because the data set has already been supplied by the pipe. We could have also written the filter as:

```
filter(wages, schooling > 10, exper > 10)
```

```
##           wage schooling      sex exper
## 1 249.6774           13 female      11
```

- 2 We didn't need to use the logical `&`. Though multiple conditions can still be written in this way with `filter()`, the default is just to separate them with a comma.



Selecting data

Just like `filter` is in many ways a more convenient form of writing out bracketed row subset conditions, the verb `select()` is largely a more convenient method for writing column arguments.

```
# Using brackets
```

```
wages_row1[,c("wage", "schooling", "exper")]
```

```
##           wage schooling exper
## 1 134.2306           13      8
```

```
# Using select
```

```
wages_row1 %>% select(wage, schooling, exper)
```

```
##           wage schooling exper
## 1 134.2306           13      8
```



An example of dropping a column

One option we have not covered so far in creating subsets is dropping rows or columns.

R has a specific notation for this, easily used with `select()`:

```
wages_row1 # What wages_row1 looks like:
```

```
##           wage schooling      sex exper
## 1 134.2306           13 female        8
```

```
wages_row1 %>% select(-exper) #drop exper
```

```
##           wage schooling      sex
## 1 134.2306           13 female
```



An example of dropping a column

Dropping columns (or rows) using the `-` notation also works with brackets, but only when using the number location of the row or column to be dropped.

```
wages_row1[, -4] # works
```

```
##           wage schooling      sex  
## 1 134.2306           13 female
```

```
# wages_row1[, -"exper"] does not work
```

Because of the ability to use name arguments, dropping with `select()` is generally easier.



“Mutating” data

Creating new variables that are functions of existing variables in a data set can be done with `mutate()`.

`mutate()` takes as its first argument the data set to be used and the equation for the new variable:

```
wages <- wages %>%  
  mutate(expsq = exper^2) # Create expsq  
wages # Display wages
```

##		wage	schooling	sex	exper	expsq
##	1	134.23058	13	female	8	64
##	2	249.67744	13	female	11	121
##	3	53.56478	10	female	11	121



Summarizing data

Summary statistics can also be easily created using the tidyverse function `summarize()`

The `summarize()` functions uses summary statistic functions in R to create a new summary tibble, with syntax largely identical to `mutate()`.

Let's try summarizing with the `mean()` summary statistic.

```
wages %>%  
  summarize(avg_wage = mean(wage))
```

```
##   avg_wage  
## 1 145.8243
```



Summary Statistics functions in R

There are a number of summary statistics available in R, which can be used either with the `summarize()` command or outside of it:

Measures of central tendency and spread:

- `mean()`, `median()`, `sd()`, `var()`, `quantile()`, `IQR()`

Position:

- `first()`, `last()`, `nth()`,

Count:

- `n()`, `n_distinct()`,



Multiple summary variables

Let's look at an example of using multiple summary variables with a larger 50-observation sample for the wages data set.

```
wages %>%  
  summarize(avg.wage = mean(wage), sd.wage = sd(wage),  
            avg.exper = mean(exper), sd.exper = sd(exper))
```

```
## # A tibble: 1 x 4  
##   avg.wage sd.wage avg.exper sd.exper  
##   <dbl>    <dbl>    <dbl>    <dbl>  
## 1     5942    17526      7.47      2.08
```



Grouping data

Creating summary statistics by group is another routine task. This is accommodated in the tidyverse using the `group_by()`.

- The arguments of `group_by()`, in addition to the data set, are simply the grouping variables separated by commas.

```
wages %>% group_by(sex) %>%  
  summarize(avg.wage = mean(wage), sd.wage = sd(wage))
```

```
## # A tibble: 2 x 3  
##   sex      avg.wage sd.wage  
##   <fct>      <dbl>   <dbl>  
## 1 female      5473    18883  
## 2 male        6410    16711
```



Arranging (sorting) data

If you want to sort your data by the values of a particular variable, you can easily do so as well with the `arrange()` function.

```
wages[1:3,] %>% arrange(exper)
```

```
## # A tibble: 3 x 4
##   wage schooling sex      exper
##   <dbl>      <int> <fct>   <int>
## 1   175         12 female     5
## 2   103         11 male       7
## 3  1411         14 female     8
```

Not: `arrange()` sorts values in ascending order by default. If you want to sort in descending order, wrap the variable name inside `desc()` in the function.



Sampling from data

Creating a sample from a data set in R is made easy by two main function in R: [sample_n](#) and [sample_frac](#).

Syntax:

- `sample_n(data, size, replace = FALSE/TRUE)`
- `sample_frac(data, size = 1, replace = FALSE/TRUE)`



A data prep example with fuel economy data

Let's use tidyverse data manipulation verbs to work through a practical data prep problem from start to finish.

For the problem, Let's use fuel economy data again, but with half of the data set. The data comes from the `vehicles` data set in the `fueleconomy` package.

```
# install.packages("fueleconomy") # Run only once  
library(fueleconomy)
```

Now let's look at how fuel efficiency has changed over time in the data set. Specifically, let's create descriptive statistics of fuel efficiency by year for "normal" passenger vehicles (4-8 cylinders).



What's in the data set?

```
glimpse(vehicles[2:12], width=50)
```

```
## Observations: 33,442
## Variables: 11
## $ make <chr> "AM General", "AM General", "AM...
## $ model <chr> "DJ Po Vehicle 2WD", "DJ Po Veh...
## $ year <int> 1984, 1984, 1984, 1984, 1985, 1...
## $ class <chr> "Special Purpose Vehicle 2WD", ...
## $ trans <chr> "Automatic 3-spd", "Automatic 3...
## $ drive <chr> "2-Wheel Drive", "2-Wheel Drive...
## $ cyl <int> 4, 4, 6, 6, 4, 6, 6, 4, 4, 6, 4...
## $ displ <dbl> 2.5, 2.5, 4.2, 4.2, 2.5, 4.2, 3...
## $ fuel <chr> "Regular", "Regular", "Regular"...
## $ hwy <int> 17, 17, 13, 13, 17, 13, 21, 26,...
## $ cty <int> 18, 18, 13, 13, 16, 13, 14, 20,...
```



View summary tibble

```
# Print annual.mpg  
annual.mpg
```

```
## # A tibble: 32 x 5  
##   year hwy.avg hwy.sd city.avg city.sd  
##   <int>   <dbl>   <dbl>   <dbl>   <dbl>  
## 1  2015    28.6    5.42    20.6    4.78  
## 2  2014    27.8    6.44    20.4    5.76  
## 3  2013    27.3    6.02    20.0    5.78  
## 4  2012    26.3    5.83    19.2    5.21  
## 5  2011    25.7    5.36    18.8    4.86  
## 6  2010    25.3    4.90    18.4    4.38  
## 7  1985    22.8    6.21    17.7    4.70  
## 8  2009    24.2    4.54    17.6    3.92  
## 9  1986    22.4    5.82    17.5    4.48
```



Cleaning data



Common data cleaning tasks

There are a few data cleaning tasks that are pervasive in empirical work:

- 1 Ensure columns have useful names
- 2 Recoding variable values
- 3 Addressing missing values



Renaming columns

Renaming columns is easily accommodated with the tidyverse `rename()` command.

- Arguments are the data set and $NewVarName = OldVarName$.

To see `rename()` in action, let's go back to the `barro.lee.data` educational data set we imported earlier:



Renaming columns example

Let's look at columns 1 and 7 through 9:

```
glimpse(barro.lee.data[,c(1,7:9)], width = 50)
```

```
## Observations: 1,898
```

```
## Variables: 4
```

```
## $ BLcode <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
```

```
## $ lu      <dbl> 80.68459, 81.05096, 82.61115, ...
```

```
## $ lp      <dbl> 17.563400, 17.018442, 14.31374...
```

```
## $ lpc     <dbl> 3.745905, 3.464397, 3.069391, ...
```

See how these variable names are uninformative?



Andrew Proctor

Renaming columns example ctd

```
barro.lee.data <- barro.lee.data %>%  
  rename(countrycode = BLcode,  
         perc.noschool = lu,  
         perc.primary = lp,  
         perc.primary.complete = lpc)
```



Renaming columns example ctd

Now let's look at the variable names again:

```
glimpse(barro.lee.data[,c(1,7:9)], width = 50)
```

```
## Observations: 1,898
```

```
## Variables: 4
```

```
## $ countrycode          <dbl> 1, 1, 1, 1, 1, ...
```

```
## $ perc.noschool        <dbl> 80.68459, 81.05...
```

```
## $ perc.primary         <dbl> 17.563400, 17.0...
```

```
## $ perc.primary.complete <dbl> 3.745905, 3.464...
```



Recoding variables

Along with renaming variables, recoding variables is another integral part of data wrangling.

```
wages[1:4,"sex"] # Look at sex column
```

```
## # A tibble: 4 x 1  
##   sex  
##   <fct>  
## 1 female  
## 2 female  
## 3 male  
## 4 male
```



Recoding variables ctd

```
wages$sex <- wages$sex %>% recode("male"=0,  
                                "female"=1) # recode  
wages[1:4,"sex"] # Look at sex column
```

```
## # A tibble: 4 x 1  
##   sex  
##   <dbl>  
## 1  1.00  
## 2  1.00  
## 3   0  
## 4   0
```



Missing Values

Another problem characteristic of observational data is missing data. In R, the way to represent missing data is with the value **NA**.

- You can recode missing value that *should be* NA but are code using a different schema either by using brackets, or the tidyverse `na_if()` function.

```
## Replace 99-denoted missing data with NA
# bracket method
wages[wages$schooling==99,] <- NA
# tidyverse method
wages$schooling <- wages$schooling %>% na_if(99)
```



Missing values continued

You can check for (correctly-coded) missing-values using the `is.na()` function.

```
## Missing  
wages[is.na(wages$wage),]
```

```
## # A tibble: 3 x 4  
##   wage schooling    sex exper  
##   <dbl>      <int> <dbl> <int>  
## 1    NA         14  1.00     8  
## 2    NA         11   0         8  
## 3    NA         10   0         8
```

Note: R does not naturally support multiple types of missingness like other languages, although it's possible to use the `sjmisc` package to do this.



Tidy data



Principles of tidy data

Rules for tidy data (from *R for Data Science*):

- ❶ Each variable must have its own column.
- ❷ Each observation must have its own row.
- ❸ Each value must have its own cell.



Tidy data tools in the tidyverse

There two main tidyverse verbs for making data tidy are:

gather(): reduces variable values are spread over multiples columns into a single column.

spread(): when multiple variables values are stored in the same columns, moves each variable into it's own column.



Gathering data

If values for a single variable are spread across multiple columns (e.g. income for different years), `gather()` moves this into single “values” column with a “key” column to identify what the different columns differentiated.

Syntax: `gather(data, key, value, columnstocombine)`



Gather example

```
earnings.panel
```

```
## # A tibble: 7 x 3
##   person y1999 y2000
##   <chr>   <dbl> <dbl>
## 1 Elsa      10.0   15.0
## 2 Mickey    20.0   28.0
## 3 Ariel     17.0   21.0
## 4 Gaston    19.0   19.0
## 5 Jasmine   32.0   35.0
## 6 Peter     22.0   29.0
## 7 Alice     11.0   15.0
```



Gather example ctd

```
earnings.panel <- earnings.panel %>%  
  gather(key="year", value="wage", y1999:y2000)  
earnings.panel
```

```
## # A tibble: 14 x 3  
##   person  year  wage  
##   <chr>   <chr> <dbl>  
## 1 Elsa    y1999    10.0  
## 2 Mickey  y1999    20.0  
## 3 Ariel   y1999    17.0  
## 4 Gaston  y1999    19.0  
## 5 Jasmine y1999    32.0  
## 6 Peter   y1999    22.0  
## 7 Alice   y1999    11.0  
## 8 Elsa    y2000    15.0
```



Spread

Spread tackles the other major problem - that often times (particularly in longitudinal data) many variables are condensed into just a “key” (or indicator) column and a value column.

```
wages2
```

##	person	indicator	values
## 1	Elsa	wage	NA
## 2	Mickey	wage	174.932480
## 3	Ariel	wage	102.668810
## 4	Gaston	wage	1.690623
## 5	Jasmine	wage	2.166231
## 6	Peter	wage	1192.371925
## 7	Alice	wage	83.363705
## 8	Elsa	wage	NA
## 9	Mickey	wage	174.932480



Spread ctd

```
wages2 %>% spread("indicator", "values")
```

##	person	wage	schooling	exper
## 1	Elsa	NA	14	8
## 2	Mickey	174.932480	12	5
## 3	Ariel	102.668810	11	7
## 4	Gaston	1.690623	11	8
## 5	Jasmine	2.166231	14	10
## 6	Peter	1192.371925	12	8
## 7	Alice	83.363705	11	6
## 8	Elsa	NA	14	8
## 9	Mickey	174.932480	12	5
## 10	Ariel	102.668810	11	7
## 11	Gaston	1.690623	11	8
## 12	Jasmine	2.166231	14	10

