

Introduction to Data Analysis in R

Module 2: Data Preparation using the Tidyverse

Andrew Proctor

andrew.proctor@phdstudent.hhs.se

January 21, 2019



- 1 Intro
- 2 Packages in R
- 3 Data Prep Preliminaries
- 4 Data Preparation
- 5 Cleaning data
- 6 Tidy data



Intro



Goals for Today

- ① Learn how to use packages in R
- ② Learn how to import and export data.
- ③ Learn how to perform common data prep functions from the tidyverse collection of packages.
- ④ Learn how to clean and “tidy” data.



Packages in R



Role of Packages in R

- Packages in R are similar to user-written commands (think *ssc install*) in Stata.
- But *most* things you do in Stata probably use core Stata commands.
- In R, most of your analysis will probably be done using packages.



[CRAN](#)
[Mirrors](#)
[What's new?](#)
[Task Views](#)
[Search](#)

[About R](#)
[R Homepage](#)

Available CRAN Packages By Date

Date	Package	
2018-02-03	fuzzyforest	Fuzzy Forests
2018-02-03	tuber	Client for the YouTube API
2018-02-02	adegenet	Exploratory Analysis of Genetic and Gen
2018-02-02	antitrust	Tools for Antitrust Practitioners
2018-02-02	arsenal	An Arsenal of 'R' Functions for Large-S

Installing and using a package

- To install a package, use the function (preferably in the console) `install.packages()`
- To begin with, let's install 2 packages:
 - `tidyverse`: the umbrella package for common data preparation and visualization in R.
 - `rio`: a package for easy data import, export (saving), and conversion.

```
install.packages("tidyverse")    # Install tidyverse  
install.packages("rio")          # Install rio
```



Loading a package during analysis

Unlike Stata, in R you need to declare what packages you will be using at the beginning of each R document.

- To do this, use the `library()` function.
 - `require()` also works, but its use is discouraged for this purpose.

```
library("tidyverse")    # Install tidyverse
library("rio")          # Install rio
```



Data Prep Preliminaries



Import and export using rio

Previously, importing and exporting data was a mess, with a lot of different functions for different file formats:

- Stata DTA files alone required two functions: `read.dta` (for Stata 6-12 DTA files), `read.dta13` (for Stata 13 and later files), etc.

The `rio` package simplifies this by reducing all of this to just one function, `import()`

- Automatically determines the file format of the file and uses the appropriate function from other packages to load in a file.



Import and export using rio II

```
PISA_2015 <- import("PISA2015.sas7bdat")  
PISA_2015[1:5,1:6]
```

##	CNTRYID	CNT	CNTSCHID	CYC	NatCen	Region
## 1	8	ALB	800001	06MS	000800	800
## 2	8	ALB	800002	06MS	000800	800
## 3	8	ALB	800003	06MS	000800	800
## 4	8	ALB	800004	06MS	000800	800
## 5	8	ALB	800005	06MS	000800	800

```
export(PISA_2015, "PISA_2015.rds")
```



Tibbles: an update to the data frame

Last class, we covered data frames—the most basic data object class for data sets with a mix of data class.

Today, we introduce one final data object: the **tibble**!

- The tibble can be thought of as an update to the data frame—and it's the first part of the *tidyverse* package that we'll look at.



Tibble vs data frames

There are three main benefits to the tibble:

① Displaying data frames:

- If you display a data frame, it will print as much as much output as allowed by the “max.print” option in the R environment. With large data sets, that’s far too much. Tibbles by default print the first 10 rows and as many columns as will fit in the window.

② Partial matching in data frames:

- When using the **\$** method to reference columns of a data frame, partial names will be matched if the reference isn’t exact. This might sound good, but the only real reason for there to be a partial match is a typo, in which case the match might be wrong.

③ Tibbles are required for some functions.



Creating or converting to tibbles

The syntax for creating tibbles exactly parallels the syntax for data frames:

- `tibble()` creates a tibble from underlying data or vectors.
- `as.tibble()` coerces an existing data object into a tibble.

```
PISA_2015 <- as.tibble(PISA_2015); PISA_2015[1:5,1:5]
```

```
## # A tibble: 5 x 5
##   CNTRYID CNT   CNTSCHID CYC   NatCen
##   <dbl> <chr>   <dbl> <chr> <chr>
## 1      8 ALB     800001 06MS 000800
## 2      8 ALB     800002 06MS 000800
## 3      8 ALB     800003 06MS 000800
## 4      8 ALB     800004 06MS 000800
## 5      8 ALB     800005 06MS 000800
```



Glimpse

Another tidyverse function that's very useful is `glimpse()`, a function very similar to `str()`.

- Both functions display information about the structure of a data object.
- `str()` provides more information, such as column (variable) attributes embedded from external data formats, but consequently is much less readable for complex data objects.
- `glimpse()` provides only column names, classes, and some data values (much more readable)
- I will often use `str()` when I want more detailed information about data structure, but use `glimpse()` for quicker glances at the data.



Pipes

Another major convenience enhancement from the tidyverse is ***pipes***, denoted **%>%**,

- Pipes allow you to combine multiple steps into a single piece of code.
- Specifically, after performing a function in one step, a pipe takes the data generated from the first step and uses it as the data input to a second step.



Pipes Example

```
barro.lee.data <- import("BL2013_MF1599_v2.1.dta") %>%  
  as.tibble() %>% glimpse(width = 50)
```

```
## Observations: 1,898  
## Variables: 20  
## $ BLcode      <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1...  
## $ country     <chr> "Algeria", "Algeria", "Al...  
## $ year        <dbl> 1950, 1955, 1960, 1965, 1...  
## $ sex         <chr> "MF", "MF", "MF", "MF", "...  
## $ agefrom     <dbl> 15, 15, 15, 15, 15, 15, 1...  
## $ age to      <dbl> 999, 999, 999, 999, 999, ...  
## $ lu          <dbl> 80.68459, 81.05096, 82.61...  
## $ lp          <dbl> 17.563400, 17.018442, 14....  
## $ lpc         <dbl> 3.745905, 3.464397, 3.069...  
## $ ls          <dbl> 1.454129, 1.639253, 2.752...
```



Data Preparation



Tidyverse and the verbs of data manipulation

A motivating principle behind the creation of the tidyverse was the language of programming should really behave like a language.

Data manipulation in the tidyverse is oriented around a few key “verbs” that perform common types of data manipulation.

- 1 `filter()` subsets the rows of a data frame based on their values.
- 2 `select()` selects variables (columns) based on their names.
- 3 `mutate()` adds new variables that are functions of existing variables.
- 4 `summarize()` creates a number of summary statistics out of many values.
- 5 `arrange()` changes the ordering of the rows.

Note: the first argument for each these functions is the data object (so pipe!).



Filtering data

Filtering keeps observations (rows) based on conditions.

- Just like using use subset conditions in the row arguments of a bracketed subset

```
# Using brackets
```

```
wages[(wages$schooling > 10) & (wages$exper > 10),]
```

```
##           wage schooling      sex exper
## 2 249.6774           13 female      11
```

```
# Using filter
```

```
wages %>% filter(schooling > 10, exper > 10)
```

```
##           wage schooling      sex exper
## 1 249.6774           13 female      11
```



Filtering data ctd

Notice a couple of things about the output:

- 1 It doesn't look like we told `filter()` what data set we would be filtering.
- That's because the data set has already been supplied by the pipe. We could have also written the filter as:

```
filter(wages, schooling > 10, exper > 10)
```

```
##           wage schooling      sex exper
## 1 249.6774          13 female      11
```

- 2 We didn't need to use the logical `&`. Though multiple conditions can still be written in this way with `filter()`, the default is just to separate them with a comma.



Selecting data

Just like `filter` is in many ways a more convenient form of writing out bracketed row subset conditions, the verb `select()` is largely a more convenient method for writing column arguments.

```
# Using brackets
```

```
wages_row1[,c("wage", "schooling", "exper")]
```

```
##           wage schooling exper
## 1 134.2306           13      8
```

```
# Using select
```

```
wages_row1 %>% select(wage, schooling, exper)
```

```
##           wage schooling exper
## 1 134.2306           13      8
```



An example of dropping a column

One option we have not covered so far in creating subsets is dropping rows or columns.

R has a specific notation for this, easily used with `select()`:

```
wages_row1 # What wages_row1 looks like:
```

```
##           wage schooling      sex exper
## 1 134.2306           13 female      8
```

```
wages_row1 %>% select(-exper) #drop exper
```

```
##           wage schooling      sex
## 1 134.2306           13 female
```



An example of dropping a column

Dropping columns (or rows) using the `-` notation also works with brackets, but only when using the number location of the row or column to be dropped.

```
wages_row1[, -4] # works
```

```
##           wage schooling    sex  
## 1 134.2306           13 female
```

```
# wages_row1[, -"exper"] does not work  
wages_row1[, "exper"] <- NULL # works (NULL is R delete)
```

Because of `select()`'s ability to use named arguments when dropping, it is generally easier (except when quotes are required due to improper names).



“Mutating” data

Creating new variables that are functions of existing variables in a data set can be done with `mutate()`.

`mutate()` takes as its first argument the data set to be used and the equation for the new variable:

```
wages <- wages %>%  
  mutate(expsq = exper^2) # Create expsq  
wages # Display wages
```

##		wage	schooling	sex	exper	expsq
## 1	134.23058		13	female	8	64
## 2	249.67744		13	female	11	121
## 3	53.56478		10	female	11	121



Summarizing data

Summary statistics can also be easily created using the tidyverse function `summarize()`

The `summarize()` functions uses summary statistic functions in R to create a new summary tibble, with syntax largely identical to `mutate()`.

Let's try summarizing with the `mean()` summary statistic.

```
wages %>%  
  summarize(avg_wage = mean(wage))
```

```
##   avg_wage  
## 1 145.8243
```



Summary Statistics functions in R

There are a number of summary statistics available in R, which can be used either with the `summarize()` command or outside of it:

Measures of central tendency and spread:

- `mean()`, `median()`, `sd()`, `var()`, `quantile()`, `IQR()`

Position:

- `first()`, `last()`, `nth()`,

Count:

- `n()`, `n_distinct()`,



Multiple summary variables

Let's look at an example of using multiple summary variables with a larger 50-observation sample for the wages data set.

```
wages %>%  
  summarize(avg.wage = mean(wage), sd.wage = sd(wage),  
            avg.exper = mean(exper), sd.exper = sd(exper))
```

```
## # A tibble: 1 x 4  
##   avg.wage sd.wage avg.exper sd.exper  
##   <dbl>    <dbl>    <dbl>    <dbl>  
## 1    5942.   17526.     7.47     2.08
```



Grouping data

Creating summary statistics by group is another routine task. This is accommodated in the tidyverse using the `group_by()`.

- The arguments of `group_by()`, in addition to the data set, are simply the grouping variables separated by commas.

```
wages %>% group_by(sex) %>%  
  summarize(avg.wage = mean(wage), sd.wage = sd(wage))
```

```
## # A tibble: 2 x 3  
##   sex      avg.wage sd.wage  
##   <fct>      <dbl>   <dbl>  
## 1 female    5473.   18883.  
## 2 male      6410.   16711.
```



Arranging (sorting) data

If you want to sort your data by the values of a particular variable, you can easily do so as well with the `arrange()` function.

```
wages[1:3,] %>% arrange(exper)
```

```
## # A tibble: 3 x 4
##   wage schooling sex      exper
##   <dbl>      <int> <fct>   <int>
## 1  175.         12 female     5
## 2  103.         11 male       7
## 3 1411.         14 female     8
```

Not: `arrange()` sorts values in ascending order by default. If you want to sort in descending order, wrap the variable name inside `desc()` in the function.



Sampling from data

Creating a sample from a data set in R is made easy by two main function in R: `sample_n` and `sample_frac`.

Syntax:

```
sample_n(data, size, replace = FALSE/TRUE)
sample_frac(data, size = 1, replace = FALSE/TRUE)
```



A data prep example with fuel economy data

Let's use tidyverse data manipulation verbs to work through a practical data prep problem from start to finish.

For the problem, Let's use fuel economy data again, but with half of the data set. The data comes from the `vehicles` data set in the `fueleconomy` package.

```
# install.packages("fueleconomy") # Run only once  
library(fueleconomy)
```

Now let's look at how fuel efficiency has changed over time in the data set. Specifically, let's create descriptive statistics of fuel efficiency by year for “normal” passenger vehicles (4-8 cylinders).



Andrew Proctor

Create summary tibble

```
annual.mpg <- vehicles %>% sample_frac(0.5) %>%  
  filter(cyl %in% 4:8) %>% group_by(year) %>%  
  summarize(hwy.avg = mean(hwy), hwy.sd = sd(hwy),  
            city.avg = mean(cty), city.sd = sd(cty)) %>%  
  arrange(desc(city.avg))
```

Note: Here I used `%in%`, which works like `inrange` in Stata. You could alternately write two inequalities to achieve the same thing.



View summary tibble

```
# Print annual.mpg  
annual.mpg
```

```
## # A tibble: 32 x 5  
##   year hwy.avg hwy.sd city.avg city.sd  
##   <int>   <dbl>   <dbl>   <dbl>   <dbl>  
## 1  2015    28.6    5.42    20.6    4.78  
## 2  2014    27.8    6.44    20.4    5.76  
## 3  2013    27.3    6.02    20.0    5.78  
## 4  2012    26.3    5.83    19.2    5.21  
## 5  2011    25.7    5.36    18.8    4.86  
## 6  2010    25.3    4.90    18.4    4.38  
## 7  1985    22.8    6.21    17.7    4.70  
## 8  2009    24.2    4.54    17.6    3.92  
## 9  1986    22.4    5.82    17.5    4.48
```



Summarizing a data set with the `summary()` function

Although the tidyverse `summarize()` function is more powerful, often you just want a quick look at summary statistics for the whole data set.

- You can easily do this with the base R `summary()` function, which produces summaries not just for data sets, but also for other R output like the results of a regression.

```
summary(wages)
```

##	wage	schooling	sex	exper
## Min. :	1.69	Min. : 8	female:15	Min. : 3
## 1st Qu.:	44.07	1st Qu.:11	male :15	1st Qu.: 6
## Median :	160.96	Median :12		Median : 7
## Mean :	5941.66	Mean :12		Mean : 7
## 3rd Qu.:	1519.01	3rd Qu.:13		3rd Qu.: 8

Cleaning data



Common data cleaning tasks

There are a few data cleaning tasks that are pervasive in empirical work:

- 1 Ensure columns have useful names
- 2 Recoding variable values
- 3 Addressing missing values



Renaming columns

Renaming columns is easily accommodated with the tidyverse `rename()` command.

- Arguments are the data set and $NewVarName = OldVarName$.

To see `rename()` in action, let's go back to the `barro.lee.data` educational data set we imported earlier:



Renaming columns example

Let's look at columns 1 and 7 through 9:

```
glimpse(barro.lee.data[,c(1,7:9)], width = 50)
```

```
## Observations: 1,898
```

```
## Variables: 4
```

```
## $ BLcode <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
```

```
## $ lu      <dbl> 80.68459, 81.05096, 82.61115, ...
```

```
## $ lp      <dbl> 17.563400, 17.018442, 14.31374...
```

```
## $ lpc     <dbl> 3.745905, 3.464397, 3.069391, ...
```

See how these variable names are uninformative?



Renaming columns example ctd

```
str(barro.lee.data[,c(1,7:9)])
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':    1898 obs. of 3 variables:
##  $ BLcode: atomic  1 1 1 1 1 1 1 1 1 1 ...
##  ..- attr(*, "label")= chr "Country Code"
##  ..- attr(*, "format.stata")= chr "%8.0g"
##  $ lu     : atomic  80.7 81.1 82.6 80.9 73.6 ...
##  ..- attr(*, "label")= chr "Percentage of No Schooling"
##  ..- attr(*, "format.stata")= chr "%7.2f"
##  $ lp     : atomic  17.6 17 14.3 14.4 19.2 ...
##  ..- attr(*, "label")= chr "Percentage of Primary"
##  ..- attr(*, "format.stata")= chr "%7.2f"
##  $ lpc    : atomic  3.75 3.46 3.07 4.01 5.23 ...
##  ..- attr(*, "label")= chr "Percentage of Primary Comp"
##  ..- attr(*, "format.stata")= chr "%7.2f"
```



Renaming columns example ctd

```
barro.lee.data <- barro.lee.data %>%  
  rename(countrycode = BLcode,  
         perc.noschool = lu,  
         perc.primary = lp,  
         perc.primary.complete = lpc)
```



Recoding variables

Along with renaming variables, recoding variables is another integral part of data wrangling.

```
wages[1:4,"sex"] # Look at sex column
```

```
## # A tibble: 4 x 1  
##   sex  
##   <fct>  
## 1 female  
## 2 female  
## 3 male  
## 4 male
```



Recoding variables ctd

```
wages$sex <- wages$sex %>% recode("male"=0,  
                                "female"=1) # recode  
wages[1:4,"sex"] # Look at sex column
```

```
## # A tibble: 4 x 1  
##       sex  
##   <dbl>  
## 1     1  
## 2     1  
## 3     0  
## 4     0
```



Missing values continued

You can check for (correctly-coded) missing-values using the `is.na()` function.

```
## Missing  
wages[is.na(wages$wage),]
```

```
## # A tibble: 3 x 4  
##   wage schooling sex exper  
##   <dbl>      <int> <dbl> <int>  
## 1    NA         14     1     8  
## 2    NA         11     0     8  
## 3    NA         10     0     8
```

Note: R does not naturally support multiple types of missingness like other languages, although it's possible to use the `sjmisc` package to do this.



Tidy data



Principles of tidy data

Rules for tidy data (from *R for Data Science*):

- ❶ Each variable must have its own column.
- ❷ Each observation must have its own row.
- ❸ Each value must have its own cell.



Tidy data tools in the tidyverse

There two main tidyverse verbs for making data tidy are:

gather(): reduces variable values are spread over multiples columns into a single column.

spread(): when multiple variables values are stored in the same columns, moves each variable into it's own column.



Gathering data

If values for a single variable are spread across multiple columns (e.g. income for different years), `gather()` moves this into single “values” column with a “key” column to identify what the different columns differentiated.

Syntax: `gather(data, key, value, columnstocombine)`



Gather example

```
earnings.panel
```

```
## # A tibble: 7 x 3
##   person y1999 y2000
##   <chr>   <dbl> <dbl>
## 1 Elsa      10      15
## 2 Mickey    20      28
## 3 Ariel     17      21
## 4 Gaston    19      19
## 5 Jasmine   32      35
## 6 Peter     22      29
## 7 Alice     11      15
```



Gather example ctd

```
earnings.panel <- earnings.panel %>%  
  gather(key="year", value="wage", y1999:y2000)  
earnings.panel
```

```
## # A tibble: 14 x 3  
##   person  year  wage  
##   <chr>   <chr> <dbl>  
## 1 Elsa    y1999    10  
## 2 Mickey  y1999    20  
## 3 Ariel   y1999    17  
## 4 Gaston  y1999    19  
## 5 Jasmine y1999    32  
## 6 Peter   y1999    22  
## 7 Alice   y1999    11  
## 8 Elsa    y2000    15
```



Spread

Spread tackles the other major problem - that often times (particularly in longitudinal data) many variables are condensed into just a “key” (or indicator) column and a value column.

```
wages2
```

##	person	indicator	values
## 1	Elsa	wage	NA
## 2	Mickey	wage	174.932480
## 3	Ariel	wage	102.668810
## 4	Gaston	wage	1.690623
## 5	Jasmine	wage	2.166231
## 6	Peter	wage	1192.371925
## 7	Alice	wage	83.363705
## 8	Elsa	wage	NA
## 9	Mickey	wage	174.932480



Spread ctd

```
wages2 %>% spread("indicator", "values")
```

##	person	wage	schooling	exper
## 1	Elsa	NA	14	8
## 2	Mickey	174.932480	12	5
## 3	Ariel	102.668810	11	7
## 4	Gaston	1.690623	11	8
## 5	Jasmine	2.166231	14	10
## 6	Peter	1192.371925	12	8
## 7	Alice	83.363705	11	6
## 8	Elsa	NA	14	8
## 9	Mickey	174.932480	12	5
## 10	Ariel	102.668810	11	7
## 11	Gaston	1.690623	11	8
## 12	Jasmine	2.166231	14	10

